Cody Lieu
10/15/2015
cal53

CS 290: Data Engineering HW 2

TenMostFrequentVisitors.scala

Completed Stages (2)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
1	takeOrdered at TenMostFrequentVisitors.scala:24 +details	2015/10/16 04:00:01	3 s	8/8			62.9 MB	
0	map at TenMostFrequentVisitors.scala:23 +details	2015/10/16 03:59:50	11 s	8/8	917.5 MB			62.9 MB

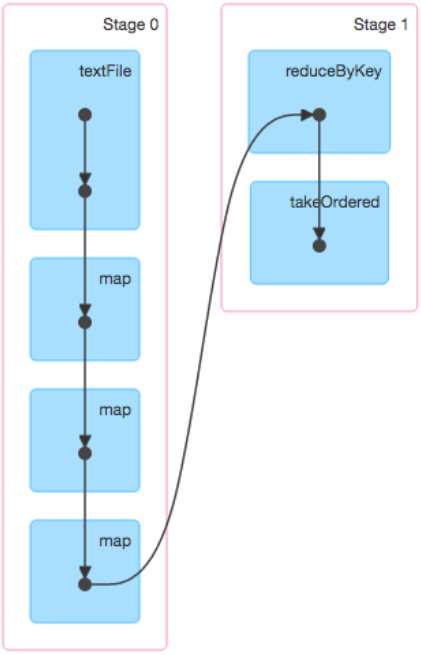

The application runs one job for computing the ten most frequent visitors. The job was comprised of 2 stages, which took a total of 14 seconds with Stage 0 taking 11 seconds and Stage 1 taking 3 seconds. Each stage had 8 tasks. 5 of the tasks in Stage 0 (0-indexed) took approximately 9 seconds to run in parallel and all ran on Executor 0 (0-indexed). All of the details of interest as well as the input and shuffle data sizes produced by each of the longest running tasks are shown below.

Tasks

Index	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Input Size / Records	Write Time	Shuffle Write Size / Records	Errors
0	0	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 03:59:51	9 s	0.7 s	128.0 MB (hadoop) / 668105	17 ms	8.4 MB / 505136	
1	1	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 03:59:51	9 s	0.7 s	128.0 MB (hadoop) / 668318	72 ms	8.8 MB / 527442	
2	2	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 03:59:51	9 s	0.7 s	128.0 MB (hadoop) / 656531	24 ms	9.6 MB / 545865	
3	3	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 03:59:51	8 s	0.7 s	128.0 MB (hadoop) / 644414	16 ms	8.8 MB / 507149	
5	5	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 03:59:51	9 s	0.7 s	128.0 MB (hadoop) / 664764	74 ms	8.0 MB / 452506	
4	4	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 03:59:51	9 s	0.7 s	128.0 MB (hadoop) / 666103	19 ms	9.1 MB / 521546	
6	6	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 03:59:51	8 s	0.7 s	128.0 MB (hadoop) / 680423	21 ms	8.8 MB / 503828	
7	7	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 03:59:51	3 s	57 ms	21.5 MB (hadoop) / 105705	6 ms	1573.5 KB / 88845	

Stage 0 was comprised of reading in the White House data file and then performing the map functions. Stage 1 was comprised of the reduceByKey call and then the takeOrdered call, which printed the final results. Looking at the bar graph below, Task 2 (0 indexed) of Stage 0 took the longest to run on Executor 0 and had an input size of 128.0 MB and a shuffle size of 9.6 MB.

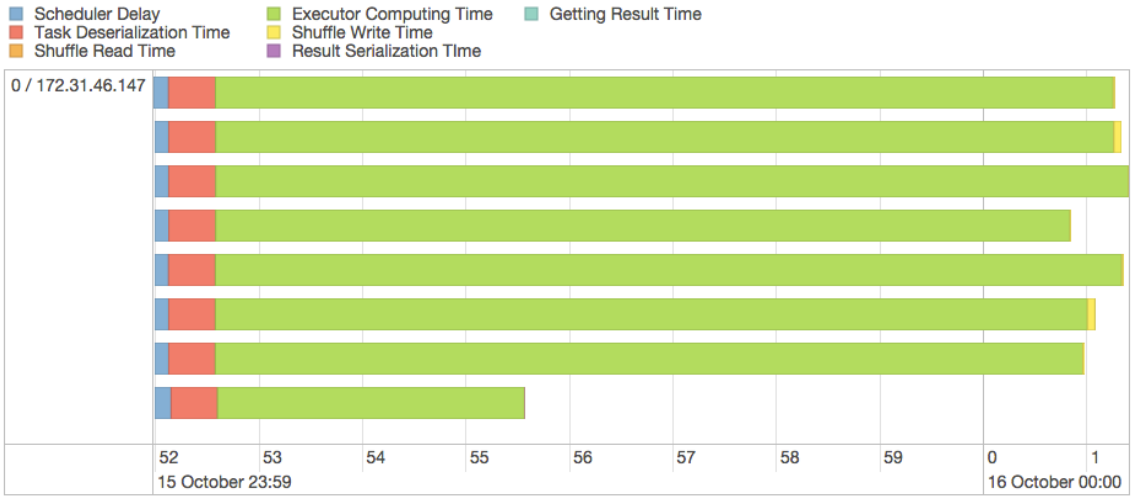

Executor data is shown below. Executor 1 completed 4 tasks and read 31.5 MB total. Executor 0 completed 12 tasks, taking in a total input of 917.5 MB and writing a total of 62.9 MB.

Details for Job 0

Status: SUCCEEDED
Completed Stages: 2

Event Timeline
Enable zooming

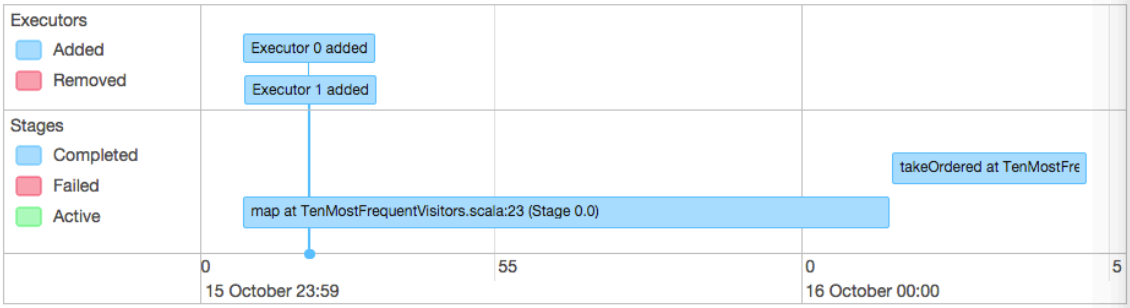

Executors (3)

Memory: 0.0 B Used (27.4 GB Total)
Disk: 0.0 B Used

Executor ID	Address	RDD Blocks	Storage Memory	Disk Used	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time	Input	Shuffle Read	Shuffle Write	Logs
0	172.31.46.147:50547	0	0.0 B / 13.6 GB	0.0 B	0	0	12	12	1.3 m	917.5 MB	0.0 B	62.9 MB	stdout stderr
1	172.31.46.148:56545	0	0.0 B / 13.6 GB	0.0 B	0	0	4	4	12.3 s	0.0 B	31.5 MB	0.0 B	stdout stderr
driver	172.31.44.186:48480	0	0.0 B / 265.4 MB	0.0 B	0	0	0	0	0 ms	0.0 B	0.0 B	0.0 B	

TenMostFrequentVisitees.scala

Completed Stages (2)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
1	takeOrdered at TenMostFrequentVisitees.scala:24 +details	2015/10/16 03:40:41	1 s	8/8			713.5 KB	
0	map at TenMostFrequentVisitees.scala:23 +details	2015/10/16 03:40:35	6 s	8/8	917.5 MB			713.5 KB

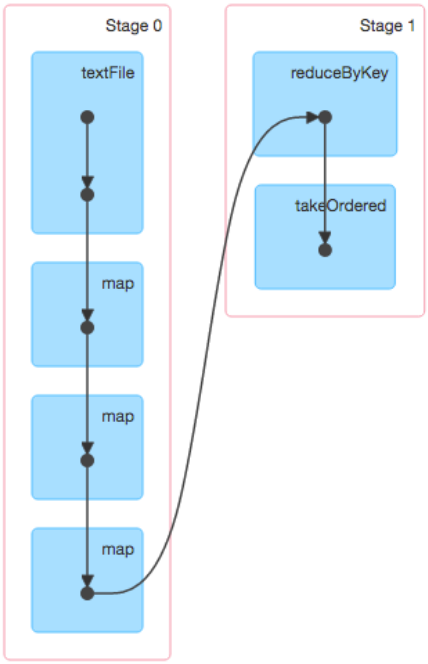

The application runs one job for computing the ten most frequent visitees. The job was comprised of 2 stages, which took a total of 7 seconds with Stage 0 taking 6 seconds and Stage 1 taking 1 second. Each stage had 8 tasks. 7 of the tasks in Stage 0 (0-indexed) took approximately 4 seconds to run in parallel and all ran on Executor 1 (0-indexed). All of the details of interest as well as the input and shuffle data sizes produced by each of the longest running tasks are shown below.

Tasks

Index	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Input Size / Records	Write Time	Shuffle Write Size / Records	Errors
0	0	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 03:40:37	4 s	55 ms	128.0 MB (hadoop) / 668105		74.8 KB / 4422	
1	1	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 03:40:37	4 s	55 ms	128.0 MB (hadoop) / 668318		73.9 KB / 4372	
2	2	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 03:40:37	4 s	55 ms	128.0 MB (hadoop) / 656531		103.6 KB / 5838	
3	3	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 03:40:37	4 s	55 ms	128.0 MB (hadoop) / 644414		88.3 KB / 4986	
5	5	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 03:40:37	4 s	55 ms	128.0 MB (hadoop) / 664764		114.5 KB / 6567	
4	4	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 03:40:37	4 s	55 ms	128.0 MB (hadoop) / 666103		95.3 KB / 5423	
6	6	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 03:40:37	4 s	55 ms	128.0 MB (hadoop) / 680423		123.2 KB / 6958	
7	7	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 03:40:37	2 s	27 ms	21.5 MB (hadoop) / 105705	1 ms	39.9 KB / 2146	

Stage 0 was comprised of reading in the White House data file and then performing the map functions. Stage 1 was comprised of the reduceByKey call and then the takeOrdered call, which printed the final results. Looking at the bar graph below, Task 6 (0 indexed) of Stage 0 took the longest to run on Executor 1 and had an input size of 128.0 MB and a shuffle size of 123.2 KB.

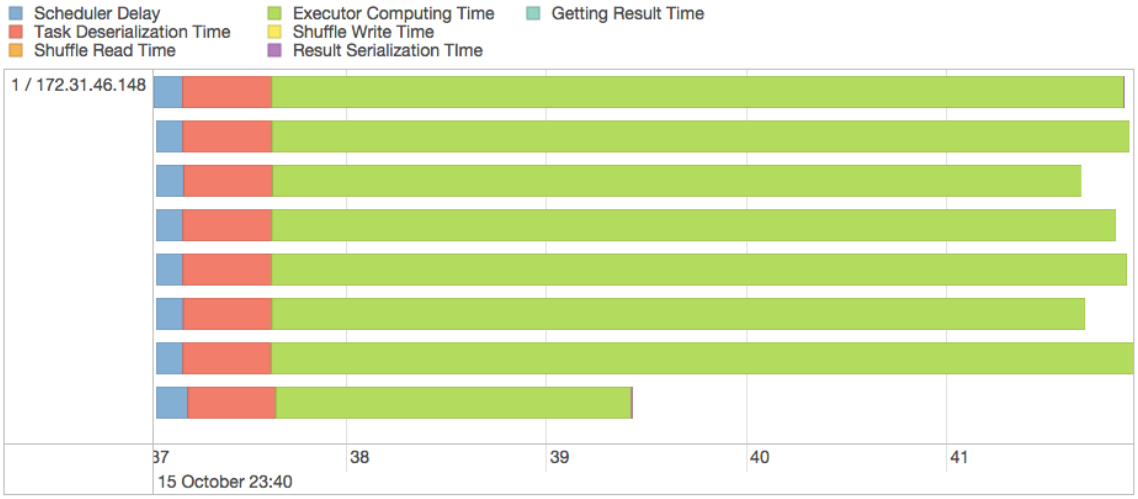

Executor data is shown below. Executor 1 completed 12 tasks, taking in a total input of 917.5 MB and writing a total of 713.5 KB. Executor 0 completed 4 tasks and read 353 KB of data total.

Details for Job 0

Status: SUCCEEDED
Completed Stages: 2

▾ Event Timeline
☐ Enable zooming

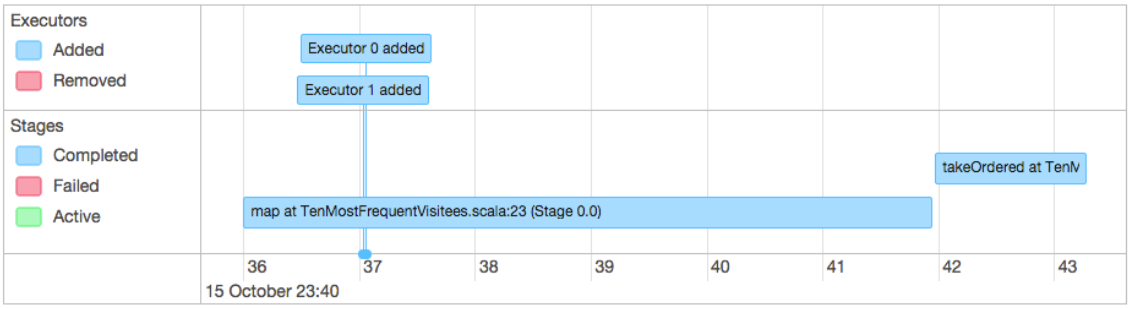

Executors (3)

Memory: 0.0 B Used (27.4 GB Total)
Disk: 0.0 B Used

Executor ID	Address	RDD Blocks	Storage Memory	Disk Used	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time	Input	Shuffle Read	Shuffle Write	Logs
0	172.31.46.147:51591	0	0.0 B / 13.6 GB	0.0 B	0	0	4	4	4.9 s	0.0 B	353.0 KB	0.0 B	stdout stderr
1	172.31.46.148:60622	0	0.0 B / 13.6 GB	0.0 B	0	0	12	12	37.2 s	917.5 MB	0.0 B	713.5 KB	stdout stderr
driver	172.31.44.186:42171	0	0.0 B / 265.4 MB	0.0 B	0	0	0	0	0 ms	0.0 B	0.0 B	0.0 B	

TenMostFrequentVisitorVisitees.scala

Completed Stages (2)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
1	takeOrdered at TenMostFrequentVisitorVisitees.scala:25 +details	2015/10/16 02:58:56	4 s	8/8			94.9 MB	
0	map at TenMostFrequentVisitorVisitees.scala:24 +details	2015/10/16 02:58:41	15 s	8/8	917.5 MB			94.9 MB

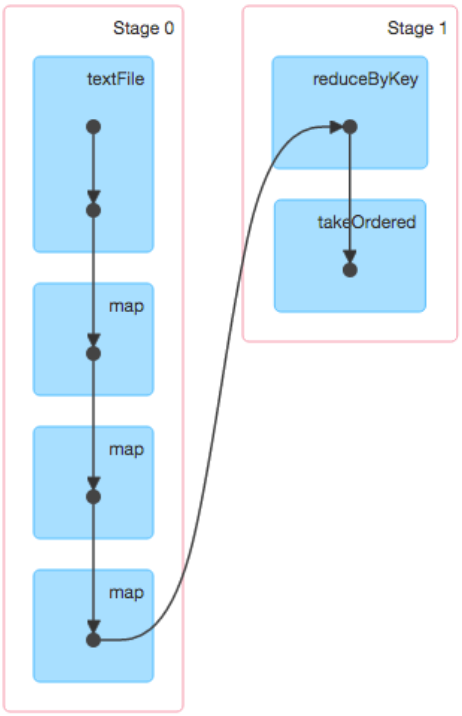

The application runs one job for computing the ten most frequent visitor-visitee combinations. The job was comprised of 2 stages, which took a total of 19 seconds with Stage 0 taking 4 seconds and Stage 1 taking 15 seconds. Each stage had 8 tasks. 7 of the tasks in Stage 0 (0-indexed) took approximately 13 seconds to run in parallel and all ran on Executor 1 (0-indexed). All of the details of interest as well as the input and shuffle data sizes produced by each of the longest running tasks are shown below.

Tasks

Index	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Input Size / Records	Write Time	Shuffle Write Size / Records	Errors
0	0	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 02:58:42	13 s	1 s	128.0 MB (hadoop) / 668105	30 ms	12.5 MB / 562473	
1	1	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 02:58:42	13 s	1 s	128.0 MB (hadoop) / 668318	22 ms	12.5 MB / 581179	
2	2	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 02:58:42	13 s	1 s	128.0 MB (hadoop) / 656531	0.1 s	13.4 MB / 590105	
3	3	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 02:58:42	13 s	1 s	128.0 MB (hadoop) / 644414	30 ms	13.0 MB / 570418	
5	5	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 02:58:42	13 s	1 s	128.0 MB (hadoop) / 664764	31 ms	13.7 MB / 552483	
4	4	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 02:58:42	13 s	1 s	128.0 MB (hadoop) / 666103	74 ms	13.4 MB / 583218	
6	6	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 02:58:42	13 s	1 s	128.0 MB (hadoop) / 680423	0.1 s	14.1 MB / 587774	
7	7	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 02:58:42	4 s	0.5 s	21.5 MB (hadoop) / 105705	8 ms	2.3 MB / 96679	

Stage 0 was comprised of reading in the White House data file and then performing the map functions. Stage 1 was comprised of the reduceByKey call and then the takeOrdered call, which printed the final results. Looking at the bar graph below, Task 6 (0 indexed) of Stage 0 took the longest to run on Executor 1 and had an input size of 128.0 MB and a shuffle size of 14.1 MB.

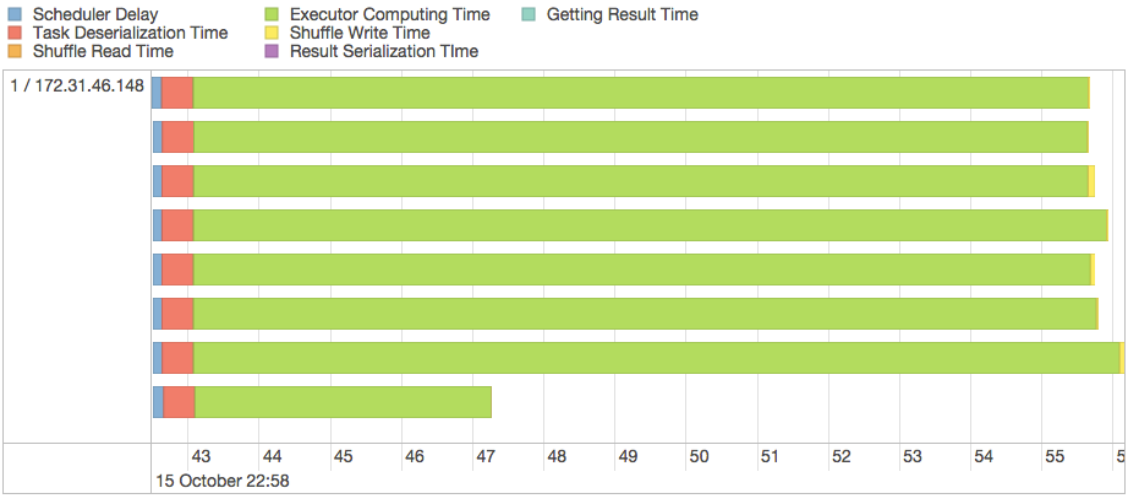

Executor data is shown below. Executor 1 completed 12 tasks, taking in a total input of 917.5 MB and writing a total of 94.9 MB. Executor 0 completed 4 tasks and read 47.4 MB of data total.

Details for Job 0

Status: SUCCEEDED
Completed Stages: 2

▾ Event Timeline
☐ Enable zooming

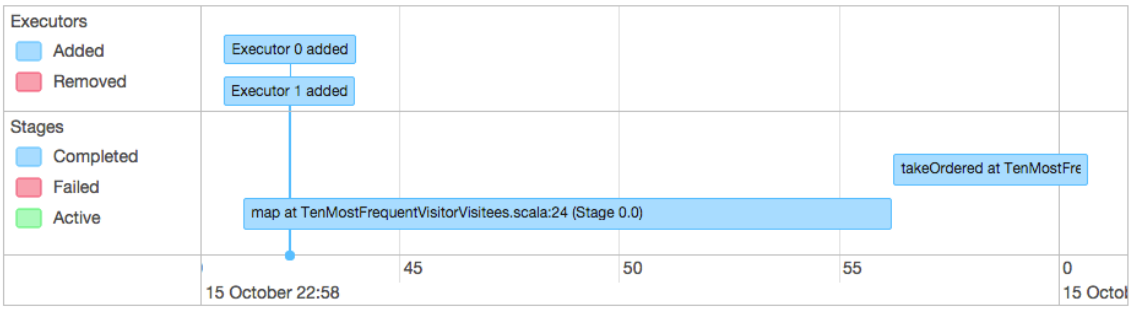

Executors (3)

Memory: 0.0 B Used (27.4 GB Total)
Disk: 0.0 B Used

Executor ID	Address	RDD Blocks	Storage Memory	Disk Used	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time	Input	Shuffle Read	Shuffle Write	Logs
0	172.31.46.147:44107	0	0.0 B / 13.6 GB	0.0 B	0	0	4	4	17.1 s	0.0 B	47.4 MB	0.0 B	stdout stderr
1	172.31.46.148:45632	0	0.0 B / 13.6 GB	0.0 B	0	0	12	12	1.8 m	917.5 MB	0.0 B	94.9 MB	stdout stderr
driver	172.31.44.186:51587	0	0.0 B / 265.4 MB	0.0 B	0	0	0	0	0 ms	0.0 B	0.0 B	0.0 B	

WikipediaPagesWithNoOutlinks.scala

Completed Jobs (2)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	count at WikipediaPagesWithNoOutlinks.scala:37	2015/10/16 04:27:06	17 s	5/5	16/16
0	zipWithIndex at WikipediaPagesWithNoOutlinks.scala:28	2015/10/16 04:27:02	3 s	1/1	1/1

The application runs two jobs with a total of 6 Stages for computing all of the Wikipedia pages with no outlinks. Job 0 was comprised of 1 stage of 1 task and took 3 seconds total. Job 1 was comprised of 5 stages and 16 tasks total, which took a total of 17 seconds.

Details for Job 0

Status: SUCCEEDED
Completed Stages: 1

▾ Event Timeline
☐ Enable zooming

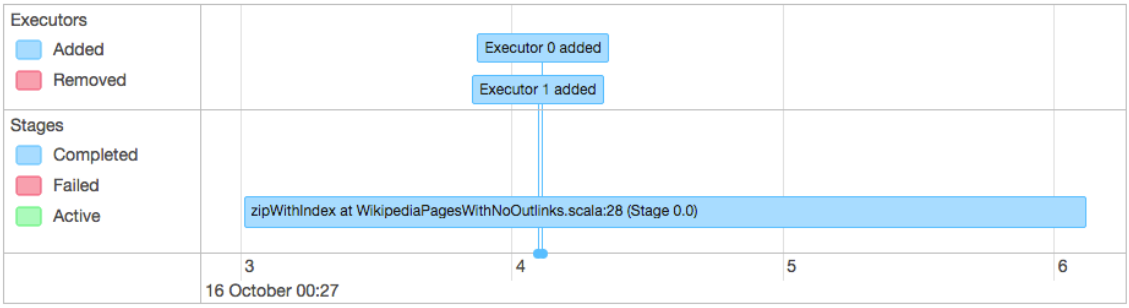

▾ DAG Visualization

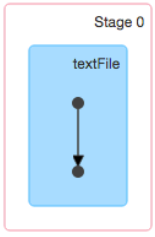

Completed Stages (1)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
0	zipWithIndex at WikipediaPagesWithNoOutlinks.scala:28 +details	2015/10/16 04:27:03	3 s	1/1	53.2 MB			

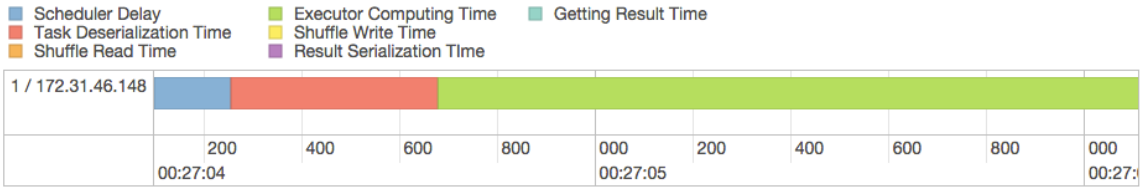

Tasks

Index	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Input Size / Records	Errors
0	0	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 04:27:04	1 s		53.2 MB (hadoop) / 2915094	

Stage 0 took 3 seconds with 1 task. Task 0 handled the zipWithIndex call and took approximately 1.5 seconds with an input size of 53.2 MB on Executor 1. This information is summarized above.

Details for Job 1

Status: SUCCEEDED
Completed Stages: 5

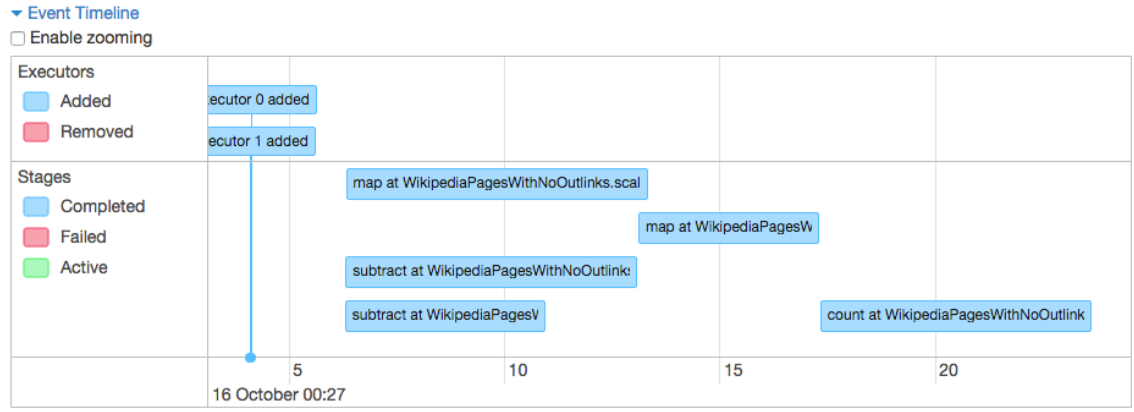

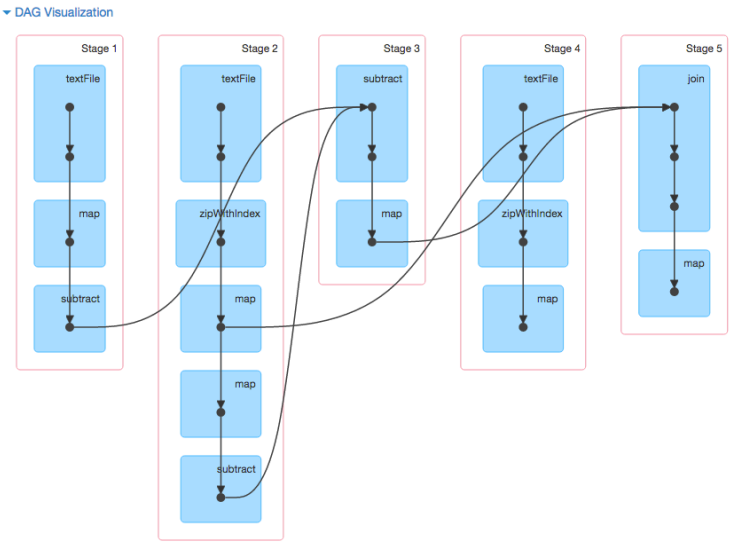

Completed Stages (5)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
5	count at WikipediaPagesWithNoOutlinks.scala:37 +details	2015/10/16 04:27:17	6 s	2/2			87.5 MB	
3	map at WikipediaPagesWithNoOutlinks.scala:35 +details	2015/10/16 04:27:13	4 s	2/2			55.6 MB	63.1 KB
4	map at WikipediaPagesWithNoOutlinks.scala:28 +details	2015/10/16 04:27:06	7 s	2/2	106.4 MB			87.5 MB
2	subtract at WikipediaPagesWithNoOutlinks.scala:33 +details	2015/10/16 04:27:06	7 s	2/2	106.4 MB			27.8 MB
1	subtract at WikipediaPagesWithNoOutlinks.scala:33 +details	2015/10/16 04:27:06	5 s	8/8	1009.4 MB			27.8 MB

Job 1 was comprised of 5 Stages. Stage 1 is comprised of 8 tasks and Stages 2 through 5 are comprised of 2 tasks. Stage 1 reads in links-simple-sorted.txt, which it maps and then performs a subtract on after getting the relevant data together in Stage 3. Stage 2 reads in titles-sorted.txt and then performs zipWithIndex, two maps, and a subtract on after getting the relevant data together in Stage 3. Stage 3 handles the subtract of data from Stage 1 and Stage 2 then maps it. Stage 4 reads in titles-sorted.txt and performs a zipWithIndex and map on it. Stage 5 takes the result of Stage 3 and one of the intermediate mapped RDDs from Stage 2 and joins them together then performs another map on it to get the final result. Tasks 0 and 1 of Stage 4 took the longest to run at 7 seconds and ran on Executors 0 and 1 respectively. Task 0 had an input size of 53.2 MB and shuffle size of 44.4 MB. Task 1 had an input size of 53.2 MB and shuffle size of 43.1 MB. This information is summarized below.

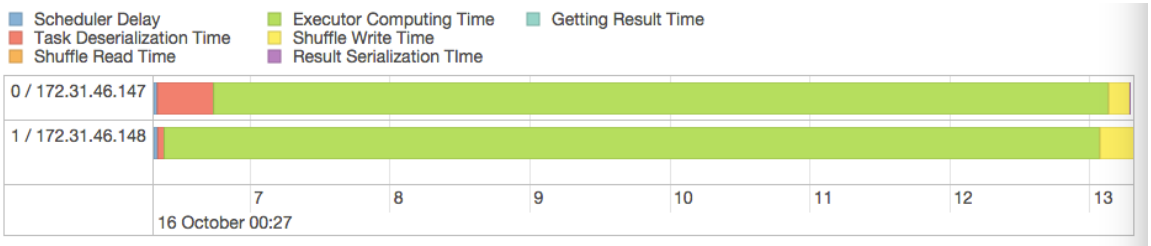

Tasks

Index	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration ▲	GC Time	Input Size / Records	Write Time	Shuffle Write Size / Records	Error
0	11	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 04:27:06	7 s	43 ms	53.2 MB (hadoop) / 2915094	0.2 s	44.4 MB / 2915094	
1	12	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 04:27:06	7 s	16 ms	53.2 MB (hadoop) /	0.3 s	43.1 MB / 2801714	

Executor data is shown below. Executor 0 completed 8 tasks, with a total input size of 603.8 MB, shuffle read size of 35.3 MB, and shuffle write size of 72.5 MB. Executor 1 completed 9 tasks, with a total input size of 671.6 MB, shuffle read size of 36.2 MB, and shuffle write size of 70.6 MB.

Executors (3)

Memory: 0.0 B Used (27.4 GB Total)
Disk: 0.0 B Used

Executor ID	Address	RDD Blocks	Storage Memory	Disk Used	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time	Input	Shuffle Read	Shuffle Write	Logs
0	172.31.46.147:41448	0	0.0 B / 13.6 GB	0.0 B	0	0	8	8	41.2 s	603.8 MB	35.3 MB	72.5 MB	stdout stderr
1	172.31.46.148:40038	0	0.0 B / 13.6 GB	0.0 B	0	0	9	9	37.8 s	671.6 MB	36.2 MB	70.6 MB	stdout stderr
driver	172.31.44.186:54247	0	0.0 B / 265.4 MB	0.0 B	0	0	0	0	0 ms	0.0 B	0.0 B	0.0 B	

WikipediaPagesWithNoInlinks.scala

Completed Jobs (2)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	count at WikipediaPagesWithNoInlinks.scala:37	2015/10/16 04:52:50	32 s	6/6	24/24
0	zipWithIndex at WikipediaPagesWithNoInlinks.scala:27	2015/10/16 04:52:47	3 s	1/1	1/1

The application runs two jobs with a total of 7 Stages for computing all of the Wikipedia pages with no inlinks. Job 0 was comprised of 1 stage of 1 task and took 3 seconds total. Job 1 was comprised of 6 stages and 24 tasks total, which took a total of 32 seconds.

Details for Job 0

Status: SUCCEEDED
Completed Stages: 1

▾ Event Timeline
☐ Enable zooming

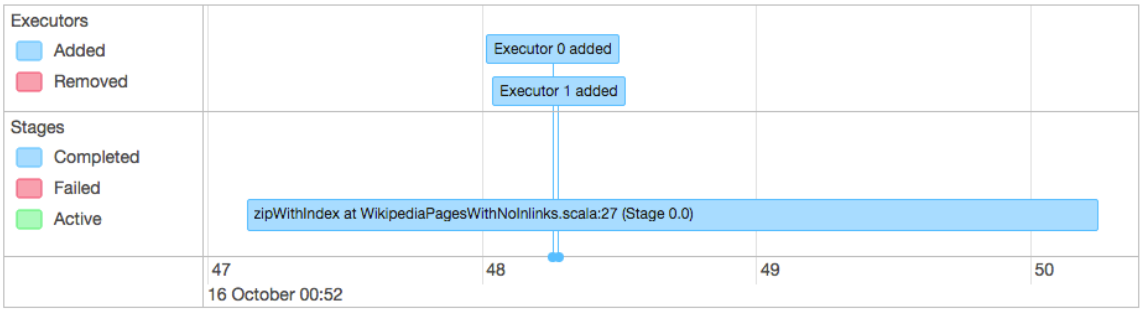

▾ DAG Visualization

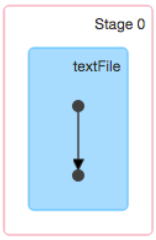

Completed Stages (1)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
0	zipWithIndex at WikipediaPagesWithNoInlinks.scala:27 +details	2015/10/16 04:52:47	3 s	1/1	53.2 MB			

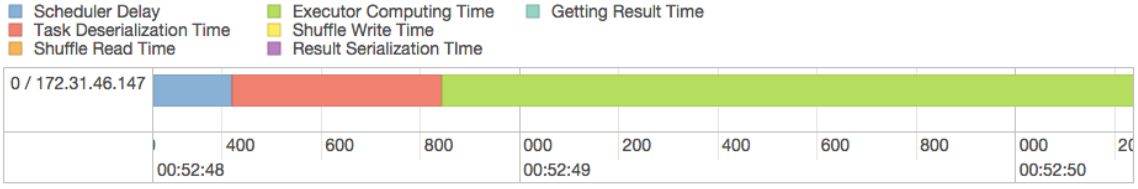

Tasks

Index	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Input Size / Records	Errors
0	0	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 04:52:48	1 s		53.2 MB (hadoop) / 2915094	

Stage 0 took 3 seconds with 1 task. Task 0 handled the zipWithIndex call and took approximately 1.5 seconds with an input size of 53.2 MB on Executor 0. This information is summarized above.

Details for Job 1

Status: SUCCEEDED
Completed Stages: 6

▾ Event Timeline
☐ Enable zooming

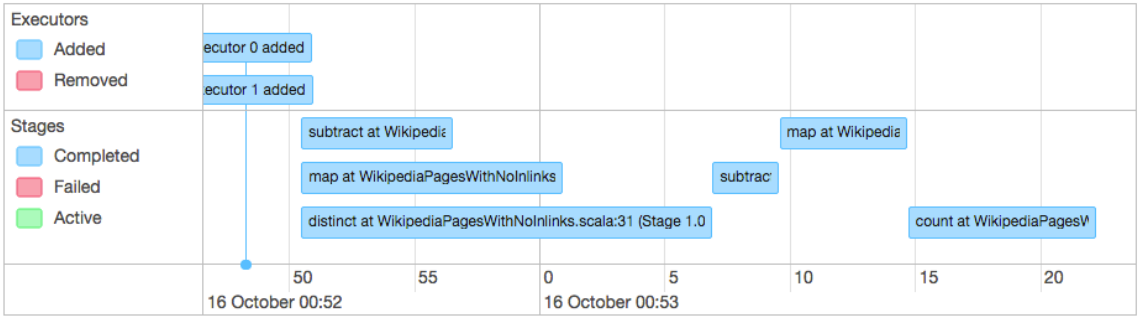

▾ DAG Visualization

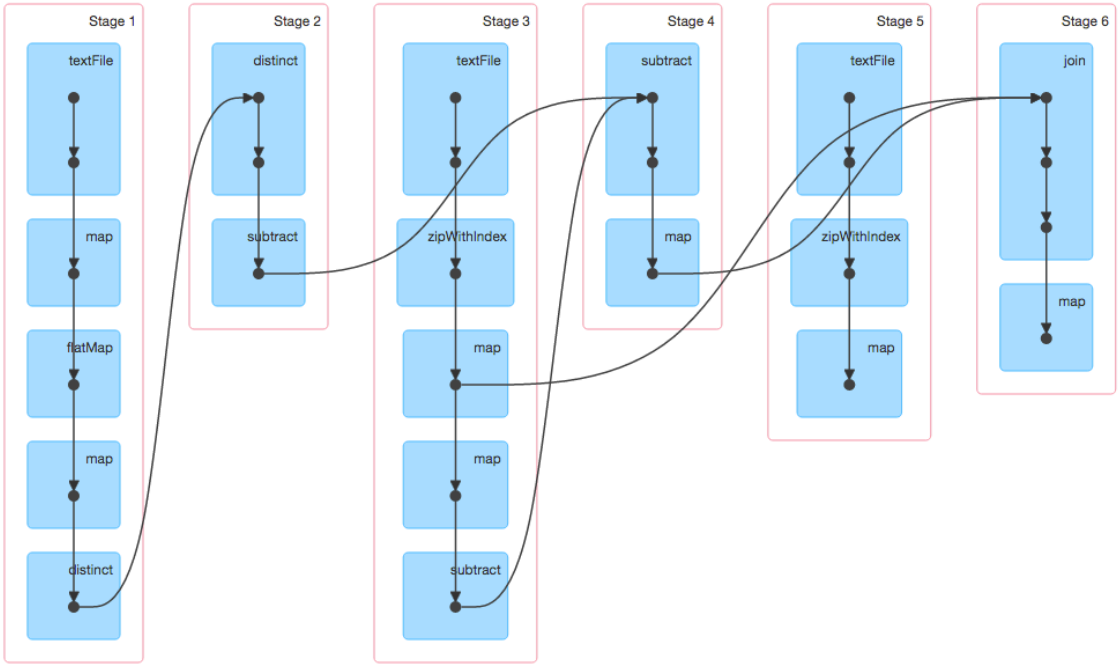

Completed Stages (6)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
6	count at WikipediaPagesWithNoInlinks.scala:37 +details	2015/10/16 04:53:14	7 s	2/2			97.5 MB	
4	map at WikipediaPagesWithNoInlinks.scala:35 +details	2015/10/16 04:53:09	5 s	2/2			49.5 MB	10.0 MB
2	subtract at WikipediaPagesWithNoInlinks.scala:33 +details	2015/10/16 04:53:06	3 s	8/8			72.7 MB	21.7 MB
5	map at WikipediaPagesWithNoInlinks.scala:27 +details	2015/10/16 04:52:50	10 s	2/2	106.4 MB			87.5 MB
3	subtract at WikipediaPagesWithNoInlinks.scala:33 +details	2015/10/16 04:52:50	6 s	2/2	106.4 MB			27.8 MB
1	distinct at WikipediaPagesWithNoInlinks.scala:31 +details	2015/10/16 04:52:50	16 s	8/8	1009.4 MB			72.7 MB

Job 1 was comprised of 6 Stages. Stages 1 and 2 are comprised of 8 tasks and Stages 3 through 6 are comprised of 2 tasks. Stage 1 reads in links-simple-sorted.txt and then performs a map, flatMap, another map, and distinct on the data. The resulting data is used in Stage 2 and subtracted with the result of Stage 3 in Stage 4. Stage 3 reads in titles-sorted.txt and then performs a zipWithIndex, two maps, then eventually a subtract in Stage 4 with the result of Stage 2. Stage 4 performs the subtract of data from Stages 2 and 3 then maps it. Stage 5 reads in titles-sorted.txt and then performs a zipWithIndex and map on it. Stage 6 takes the result from Stage 4 and an intermediate mapped RDD from Stage 3 and joins them together then performs a final map to get the resulting RDD. All of the longest tasks are in Stage 1, with 5 tasks taking approximately 16 seconds to complete. Task 0 of Stage 1 took the longest to run on Executor 0 by fractions of a second. Task 0 had an input size of 128 MB and a shuffle size of 8.5 MB. This information is summarized below.

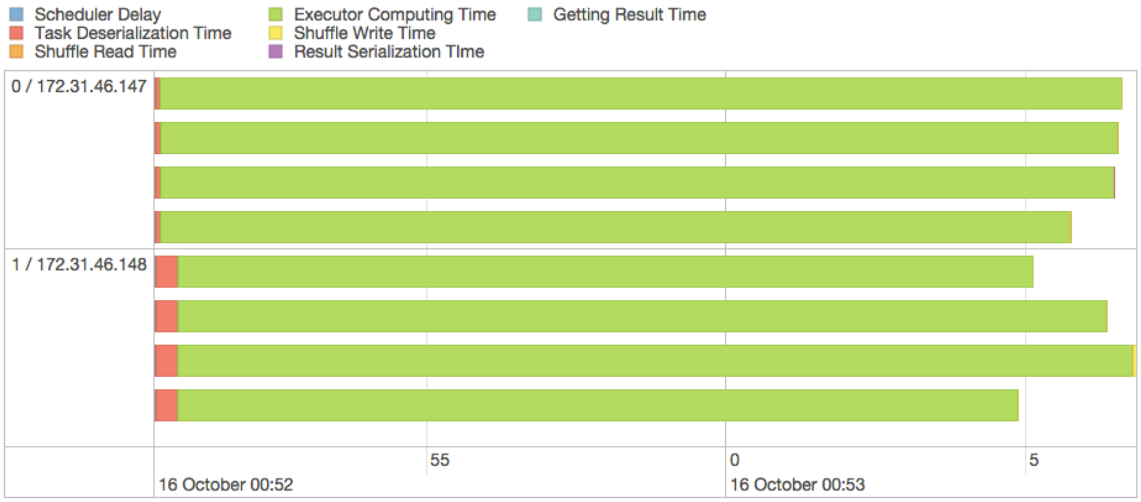

Tasks

Index	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration ▾	GC Time	Input Size / Records	Write Time	Shuffle Write Size / Records	Errors
0	1	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 04:52:50	16 s	0.5 s	128.0 MB (hadoop) / 634912	15 ms	8.5 MB / 1740228	
5	6	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 04:52:50	16 s	0.6 s	128.0 MB (hadoop) / 787260	90 ms	9.0 MB / 1845995	
2	3	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 04:52:50	16 s	0.5 s	128.0 MB (hadoop) / 813206	15 ms	9.0 MB / 1850944	
4	5	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 04:52:50	16 s	0.5 s	128.0 MB (hadoop) / 522056	19 ms	11.3 MB / 2324823	
3	4	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 04:52:50	16 s	0.6 s	128.0 MB (hadoop) / 759860	14 ms	8.9 MB / 1825635	
6	7	0	SUCCESS	ANY	0 / 172.31.46.147	2015/10/16 04:52:50	15 s	0.5 s	128.0 MB (hadoop) / 775011	16 ms	8.9 MB / 1827668	
1	2	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 04:52:50	14 s	0.6 s	128.0 MB (hadoop) / 741608	14 ms	8.9 MB / 1819931	
7	8	0	SUCCESS	ANY	1 / 172.31.46.148	2015/10/16 04:52:50	14 s	0.6 s	113.4 MB (hadoop) / 672157	14 ms	8.3 MB / 1694130	

Executor data is shown below. Executor 0 completed 13 tasks, with a total input size of 671.6 MB, shuffle read size of 61.7 MB, and shuffle write size of 112.2 MB. Executor 1 completed 12 tasks, with a total input size of 603.8 MB, shuffle read size of 64 MB, and shuffle write size of 107.5 MB.

Executors (3)

Memory: 0.0 B Used (27.4 GB Total)
Disk: 0.0 B Used

Executor ID	Address	RDD Blocks	Storage Memory	Disk Used	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time	Input	Shuffle Read	Shuffle Write	Logs
0	172.31.46.147:56909	0	0.0 B / 13.6 GB	0.0 B	0	0	13	13	1.7 m	671.6 MB	61.7 MB	112.2 MB	stdout stderr
1	172.31.46.148:59739	0	0.0 B / 13.6 GB	0.0 B	0	0	12	12	1.7 m	603.8 MB	64.0 MB	107.5 MB	stdout stderr
driver	172.31.44.186:57431	0	0.0 B / 265.4 MB	0.0 B	0	0	0	0	0 ms	0.0 B	0.0 B	0.0 B	